



Sistemas de Operação / Fundamentos de Sistemas Operativos

Memory management

Artur Pereira <artur@ua.pt>

DETI / Universidade de Aveiro

Contents

- 1 Introduction
- 2 Address space of a process
- 3 Real memory organization
- 4 Memory partitioning
- 5 Virtual memory organization
- 6 Paging
- 7 Segmentation
- 8 Page replacement
- 9 Bibliography

Memory management

Introduction

- To be executed, a process must have its address space, at least partially, resident in main memory
- In a multiprogrammed environment, to maximize processor utilization and improve response time (or turnaround time), a computer system must maintain the address spaces multiple processes resident in main memory
- But, there may not be room for all
 - because, although the main memory has been growing over the years, it is a fact that “data expands to fill the space available for storage”

(Corollary of the Parkinson's law)

Memory management

Memory hierarchy

- Ideally, an application programmer would like to have infinitely large, infinitely fast, non-volatile and inexpensive available memory
 - In practice, this is not possible
- Thus, the memory of a computer system is typically organized at different levels, forming a hierarchy
 - **cache memory** – small (tens of KB to some MB), very fast, volatile and expensive
 - **main memory** – medium size (hundreds of MB to hundreds of GB), volatile and average access speed and price
 - **secondary memory** – large (tens, hundreds or thousands of GB), slow, non-volatile and cheap

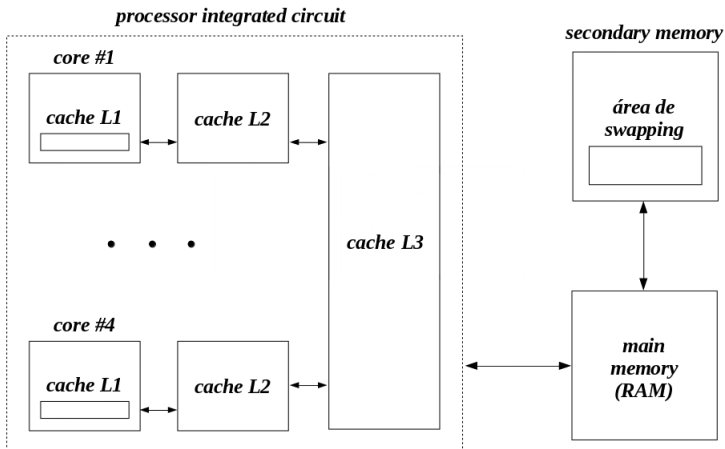
Memory management

Memory hierarchy (2)

- The **cache memory** will contain a copy of the memory positions (instructions and operands) most frequently referenced by the processor in the near past
 - The cache memory is located on the processor's own integrated circuit (**level 1**)
 - And on an autonomous integrated circuit glued to the same substrate (**levels 2 and 3**)
 - Transferring data to and from main memory is done almost completely transparent to the system programmer
- **Secondary memory** has two main functions
 - **File system** – storage for more or less permanent information (programs and data)
 - **Swapping area** – Extension of the main memory so that its size does not constitute a limiting factor to the number of processes that may currently exist.


Memory management

Memory hierarchy (3)



Memory management

Memory hierarchy (4)



- This type of organization is based on the assumption that the further an instruction or operand is away from the processor, the less times it will be referenced
 - In these conditions, the average time for a reference tends to approach the lowest value
- Based on the **principle of locality of reference**
 - The tendency of a program to access the same set of memory locations repetitively over a short period of time

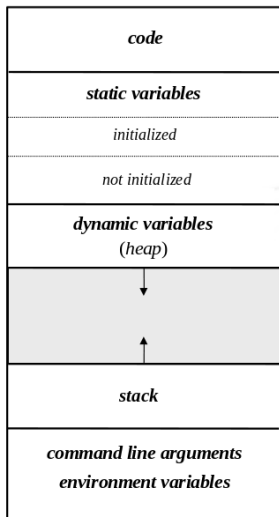
Memory management

Role

- The role of memory management in a multiprogramming environment focuses on controlling the transfer of data between main and secondary memory (**swapping area**), in order to
 - **Maintaining a register** of the parts of the main memory that are occupied and those that are free
 - **Reserving portions** of main memory for the processes that will need it, or **releasing** them when they are no longer needed
 - **Swapping out** all or part of the address space of a process when the main memory is too small to contain all the processes that coexist.
 - **Swapping in** all or part of the address space of a process when the main memory is available

Address space

Address space of a process



- **Code** and **static variables** regions have a fixed size, which is determined by the loader
- **Dynamic variables** and **stack** regions growth (in opposite directions) during the execution of the process
- It is a common practice to leave an unallocated memory area in the process address space between the dynamic definition region and the stack that can be used alternatively by any of them
- When this area is exhausted on the stack side, the execution of the process cannot continue, resulting in the occurrence of a fatal error: **stack overflow**

Address space


Address space of a process

- The binary image of the process address space represents a relocatable address space, the so-called **logical address space**
- The main memory region where it is loaded for execution, constitutes the **physical address space** of the process
- Separation between the logical and physical address spaces is a central concept to the memory management mechanisms in a multiprogrammed environment
- There are two issues that have to be solved
 - **dynamic mapping** – ability to convert a logical address to a physical address at runtime, so that the physical address space of a process can be placed in any region of main memory and be moved if necessary
 - **dynamic protection** – ability to prevent at runtime access to addresses located outside the process's own address space

Real memory organization

Logical and physical address spaces


- In a real memory organization, there is a **one-to-one correspondence** between the **logical address space** of a process and its **physical address space**



Real memory organization

Logical and physical address spaces

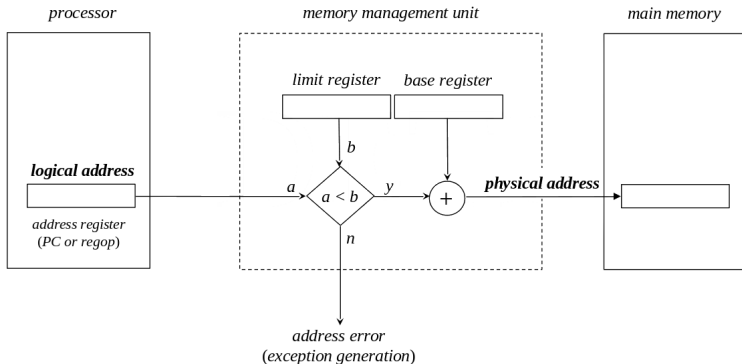
- Consequences:
 - **Limitation of the address space of a process** – in no case can memory management support automatic mechanisms that allow the address space of a process to be larger than the size of the main memory available
 - The use of overlays can allow to overcome that
 - **Contiguity of the physical address space** – although it is not a strictly necessary condition, it is naturally simpler and more efficient to assume that the process address space is contiguous
 - **Swapping area as an extension of the main memory** – it serves as storage for the storage of the address space of processes that cannot be directly loaded into main memory due to lack of space



Real memory organization

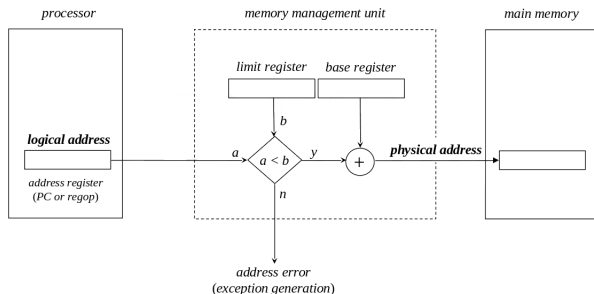
Logical address to physical address translation

- How are **dynamic mapping** and **dynamic protection** accomplished?
- A piece of hardware (the MMU) is required



Real memory organization

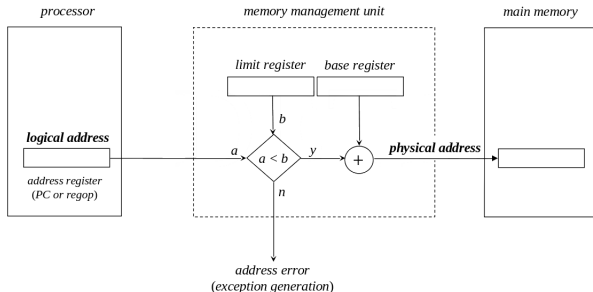
Logical address to physical address translation (2)



- The **limit register** must contain the size in bytes of the logical address space
- The **base register** must contain the address of the beginning of the main memory region where the physical address space of the process is placed
- On process switching, the **dispatch** operation loads the base and limit registers with the values present in the corresponding fields of the process control table entry associated with the process that being scheduled for execution

Real memory organization

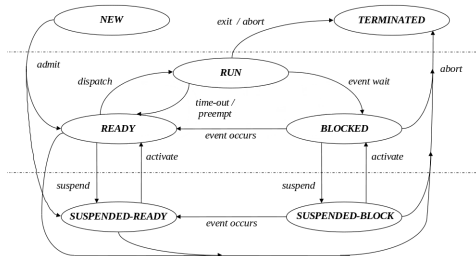
Logical address to physical address translation (2)



- Whenever there is a reference to memory
 - the logical address is first compared to the value of the limit register
 - if it is less, it is a valid reference (it occurs within the process address space) then the logical address is added to the value of the base register to produce the physical address
 - if it is greater than or equal it is an invalid reference, then a null memory access (dummy cycle) is set in motion and an exception is generated due to address error

Real memory organization


Long-term scheduling



- When a process is created, the data structures to manage it is initialized
 - Its logical address space is constructed, and the value of the limit register is computed and saved in the corresponding field of the process control table (PCT)
- If there is space in main memory, its address space is loaded there, the base register field is updated with the initial address of the reserved region and the process is placed in the READY queue
- Otherwise, its address space is temporarily stored in the swapping area and the process is placed in the SUSPENDED-READY queue

Real memory organization


Medium-term scheduling



- If memory is required to another process, a BLOCKED (or even READY) process may be **swapped out**, freeing the physical memory it is using,
 - In such a case, its base register field in the PCT becomes undefined
- If memory becomes available, a SUSPENDED-READY (or even SUSPENDED-BLOCKED) process may be **swapped in**,
 - Its base register field in the PCT is updated with its new physical location
 - A SUSPENDED-BLOCK process is only selected if no SUSPENDED-READY one exists
- When a process terminates, it is **swapped out** (if not already there), waiting for the end of operations

Memory partitioning


Fixed partitioning



- Main memory can be divided into a number of static partitions at system generation time
 - not necessarily all the same size
- The logical address space of a process may be loaded into a partition of equal or greater size
 - thus, the largest partition determines the size of the largest allowable process
- Some features:
 - Simple to implement
 - Efficient – little operating system overhead
 - Maximum number of active processes is fixed
 - Inefficient use of memory due to internal fragmentation – the part of a partition not used by a process is wasted

Memory partitioning

Fixed partitioning (2)




- If a partition becomes available, which of the SUSPENDED-READY processes should be placed there?
- Two different scheduling policies are here considered
 - **Valuing fairness** – the first process in the queue of SUSPENDED-READY processes whose address space fits in the partition is chosen
 - **Valuing the occupation of main memory** – the first process in the queue of SUSPENDED-READY processes with the largest address space that fits in the partition is chosen
 - to avoid starvation an aging mechanism can be used

Memory partitioning

Dynamic partitioning

- In dynamic partitioning, at start, all the available part of the memory constitutes a single block and then
 - reserve a region of sufficient size to load the address space of the processes that arises
 - release that region when it is no longer needed




Memory partitioning

Dynamic partitioning (2)

- As the memory is dynamically reserved, the operating system has to keep an updated record of occupied and free regions
- One way to do this is by building two (bi)linked lists
 - **list of occupied regions** – locates the regions that have been reserved for storage of the address spaces of processes resident in main memory
 - **list of free regions** – locates the regions still available
- Memory is not allocated in byte boundaries, because
 - useless very small free regions may appear
 - that will be included in the list of free regions
 - making subsequent searches more complex
- Thus, the main memory is typically divided into blocks of fixed size and allocation is made in units of these blocks


Memory partitioning

Dynamic partitioning (3)




	size (bytes)	size (BS)
main memory	256 M	65536
operating system	4M	1024
block size (BS)	4K	1
process A	4M	1024
Process C	3M	768
Process D	1M	256

lista of occupied regions



lista of free regions



Memory partitioning

Dynamic partitioning (4)

- **Valuing fairness** is the scheduling discipline generally adopted, being chosen the first process in the queue of SUSPENDED-READY processes whose address space can be placed in main memory
- Dynamic partitioning can produce **external fragmentation**
 - Free space is splitted in a large number of (possible) small free regions
 - Situations can be reached where, although there is enough free memory, it is not continuous and the storage of the address space of a new process is no longer possible
- The solution is **garbage collection** – compact the free space, grouping all the free regions into a single one
 - This operation requires stopping all processing and, if the memory is large, it can have a very long execution time

Memory partitioning

Dynamic partitioning (5)

- What free region to use to allocate the address space of a process?
- Possible policies:
 - **first fit** – the list of free regions is searched from the beginning until the first region with sufficient size is found
 - **next fit** – is a variant of the first fit which consists of starting the search from the stop point in the previous search
 - **best fit** – the list of free regions is fully researched, choosing the smallest region with sufficient size for the process
 - **worst fit** – the list of free regions is fully researched, choosing the largest existing region
- Which one is the best?
 - in terms of fragmentation
 - in terms of efficiency of allocation
 - in terms of efficiency of release

Memory partitioning

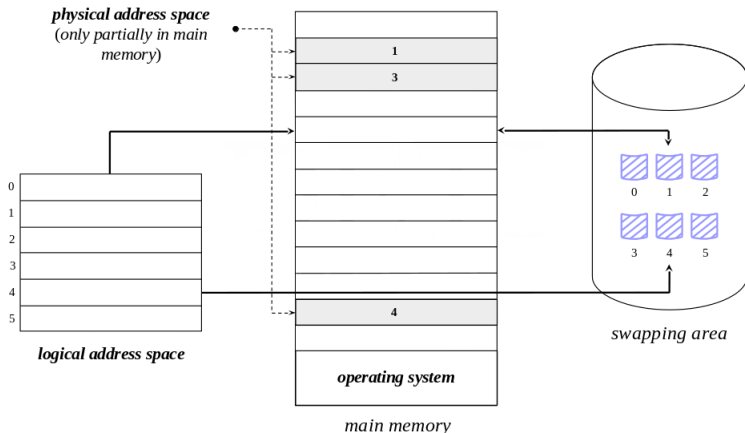
Dynamic partitioning (6)

- Advantages
 - **general** – the scope of application is independent of the type of processes that will be executed
 - **low complexity implementation** – no special hardware required and data structures are reduced to two (bi)linked lists
- Disadvantages
 - **external fragmentation** – the fraction of the main memory that ends up being wasted, given the small size of the regions in which it is divided, can reach in some cases about a third of the total (50% rule)
 - **inefficient** – it is not possible to build algorithms that are simultaneously very efficient in allocating and freeing space

Virtual memory organization

Logical and physical address spaces



- In a virtual memory organization, the **logical address space** of a process and its **physical address space** are **totally dissociated**
 - even, partially in memory and partially only in the swapping area



Virtual memory organization

Logical and physical address spaces

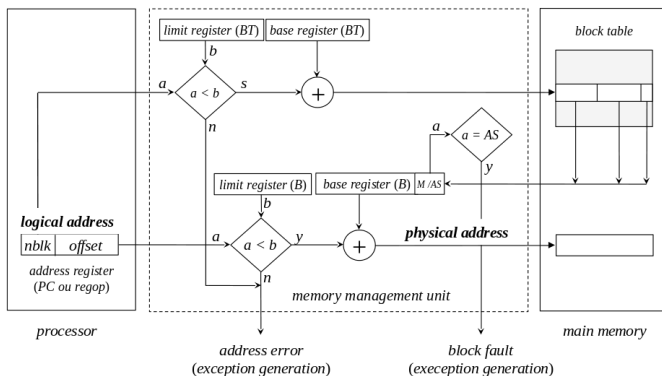
- Consequences:
 - No limitation of the address space of a process** – methodologies allowing the execution of processes whose address spaces are greater than the size of the available main memory can be established
 - Non-contiguity of the physical address space** – the address spaces of the processes, divided into blocks of fixed or variable size, are dispersed throughout the memory, trying to guarantee a more efficient occupation of the available space
 - Swapping area as an extension of the main memory** – its role is to maintain an updated image of the address spaces of the processes that currently coexist, namely their variable part (static and dynamic definition areas and stack)



Virtual memory organization

Logical address to physical address translation

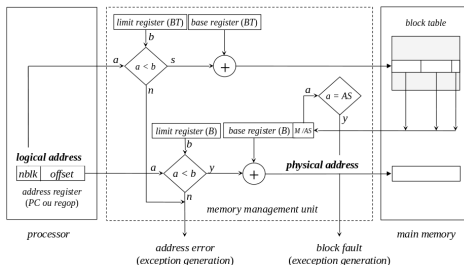
- How are **dynamic mapping** and **dynamic protection** accomplished?



- The **logical address** is composed of two fields:
 - nblk** – that identifies a specific block
 - offset** – that identifies a position within the block, as an offset from its beginning

Virtual memory organization

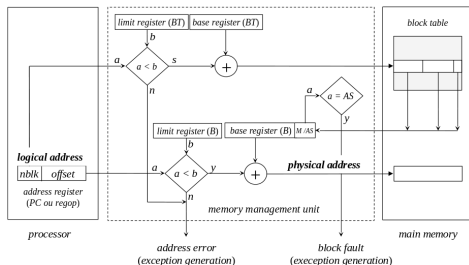
Logical address to physical address translation (2)



- The MMU contains two pairs of base and limit registers
 - one is related to the beginning and size of the **block table**, data structure describing the several blocks the address space is divided in
 - the other describes a specific block, the one being accessed
- On process switching, the **dispatch** operation loads those related to the block table, with the values stored in the PCT entry of the process scheduled for execution
 - the **base register** represents the start address of process's block table
 - the **limit register** represents the number of table entries (number of blocks)

Virtual memory organization

Logical address to physical address translation (3)



• A memory access unfolds into three steps

• Step 1:

- field **nblk** of the logical address is compared with **block table limit register**
- if it is valid (\leq), the **block table base register** plus **nblk** points to the **block table entry**, which is loaded into the MMU
- if not ($>$), a null memory access (dummy cycle) is set in motion and an **exception** is generated due to **address error**

• Step 2:

- flag **M/AS** is evaluated
- if it is **M** (the referenced block is in memory), the operation may proceed
- if not (the referenced block is swapped out), a null memory access (dummy cycle) is set in motion and an **exception** is generated due to **block fault**

• Step 3:

- field **offset** of the logical address is compared with **block limit register**
- if it is valid (\leq), the **block base register** plus **offset** points to the **physical address**
- if not ($>$), a null memory access (dummy cycle) is set in motion and an **exception** is generated due to **address error**

Virtual memory organization

Analysis

- The increase in versatility, introduced by a virtual memory organization, has the cost of transforming each memory access into two accesses
 - in the first, the process block table is accessed, to obtain the address of the beginning of the block in memory
 - only in the second the specific memory position is accessed
- Conceptually, the virtual memory organization results in a partitioning of the logical address space of the process in blocks that are dynamically treated as autonomous address sub-spaces in a real memory organization
 - of fixed partitions, if the blocks are all the same size
 - of variable partitions, if they can have a different size
- What's new is the possibility of access to a block currently not resident in main memory
 - with the consequent need to cancel the access and repeat it later, when the block is loaded

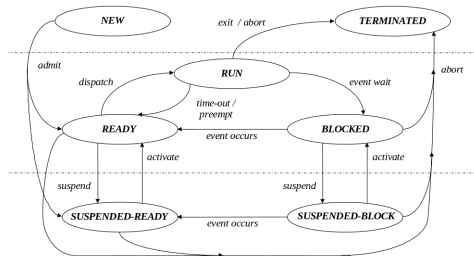
Virtual memory organization

Role of the TLB

- The need for this double access to memory can be minimized by taking advantage of the **principle of locality of reference**
- As the accesses will tend to be concentrated in a well-defined set of blocks during extended process execution time intervals, the memory management unit usually keeps the content of the block table entries stored in an internal associative memory, called the **translation lookaside buffer (TLB)**
- Thus, the first access can be a
 - **hit** – when the entry is stored in the TLB, in which case the access is internal to the processor
 - **miss** – when the entry is not stored in the TLB, in which case there is access to the main memory
- The average access time to an instruction or operand tends to approximate the lowest value
 - an access to the TLB plus an access to the main memory

Virtual memory organization


Long-term scheduling



- When a process is created, the data structures to manage it is initialized
 - Its logical address space is constructed, at least its variable part is put in the swapping area and its block table is organized
 - Some blocks can be shared with other processes
- If there is space in main memory, at least its block table, first block of code and block of the its stack are loaded there, the corresponding entries in the block table are updated and the process is placed in the READY queue
- Otherwise, the process is placed in the SUSPENDED-READY queue

Virtual memory organization


Short-term scheduling



- During its execution, on **block fault**, a process is placed in the **BLOCKED** state, while the **faulty block is swapped in**
- When the block is in memory, the process is placed in the **READY** state

Virtual memory organization

Medium-term scheduling



- While READY or BLOCKED, all blocks of a process may be **swapped out**, if memory space is required
- While SUSPENDED-READY (or SUSPENDED-BLOCKED), if memory space becomes available, the block table and a selection of blocks of a process may be **swapped in**, and the process transitions to READY or BLOCKED
 - the corresponding entries of the block table are updated
 - A SUSPENDED-BLOCK process is only selected if no SUSPENDED-READY one exists
- When a process terminates, it is **swapped out** (if not already there), waiting for the end of operations

Virtual memory organization

Block fault exception

- A relevant property of the virtual memory organization is the capacity of the computer system for executing processes whose address space is not in its entirety, and simultaneously, residing in main memory
- The operating system has to provide means to solve the problem of referencing an address located in a block that is not currently present in memory
- The MMU generates an exception in these circumstances and its service routine must start actions to **swap in** the block and, after its completion, **repeat** the execution of the instruction that produced the fault
- All these operations are carried out in a **completely transparent** way to the user who is not aware of the interruptions introduced in the execution of the process

Virtual memory organization

Block fault exception – service procedure

- The context of the process is saved in the corresponding entry of the process control table, its state changes to BLOCKED and the program counter is updated to the address that produced the block fault
- If memory space is available, a region is selected to swap in the missing block
- If not, an occupied block is selected to be replaced
 - if this block was modified, it is transferred to the swapping area
 - its entry in the block table is updated to indicate it is no longer in memory
- The swapping in of the missing block is started (to the selected region)
- The dispatch function of the processor scheduler is called to put in execution one of the READY process
- When the transfer is concluded, the block table entry is updated and the process is put in the READY state


Virtual memory organization

Block fault exception – analysis

- If, during its execution, a process was continuously generating block fault exceptions, the processing rate would be very slow and, in general, the throughput of the computer system would also be very low, jeopardizing the usefulness of a organization of virtual memory in multiprogramming
- In practice, this is not the case, because of the principle of **locality of reference**
 - A process, whatever the stage of its execution, accesses during a relatively extended periods of time only a small fraction of its address space
 - This fraction naturally changes over time, but in each interval considered there are typically thousands of accesses that are made on well-defined fractions of its address space.
 - As long as the corresponding blocks are in memory, the rhythm of execution can proceed without the occurrence of block faults

Paging

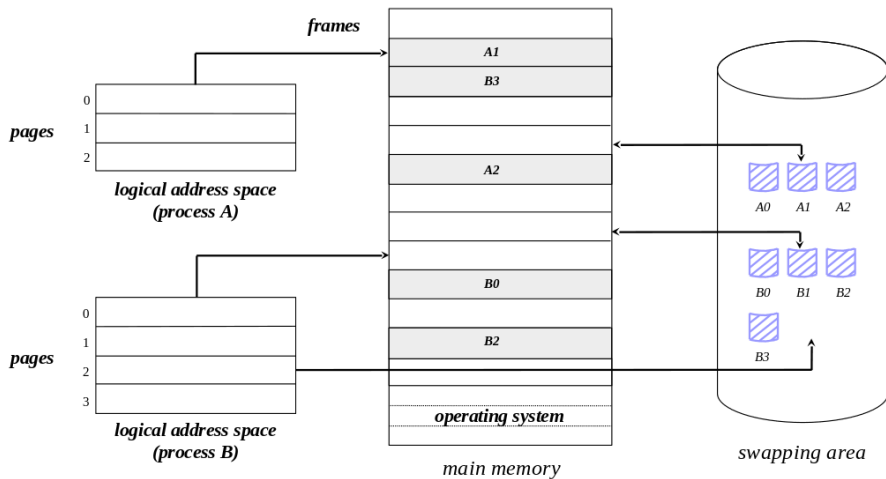
Introduction



- Memory is divided into equal fixed-size chunks, called **frames**
 - a power of 2 is used for the size, typically 4 or 8 KB
- The address space of a process is divided into fixed-size blocks, of the same size, called **pages**
- While dividing the address space into pages, the linker usually starts a new page when a new region starts
- In a logical address:
 - the most significant bits represent the **page number**
 - the least significant bits represent an **offset** within the page

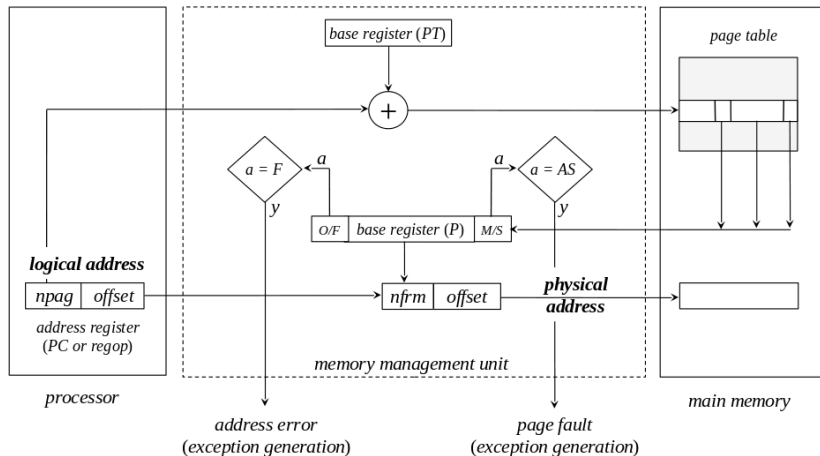
Paging

Illustration example



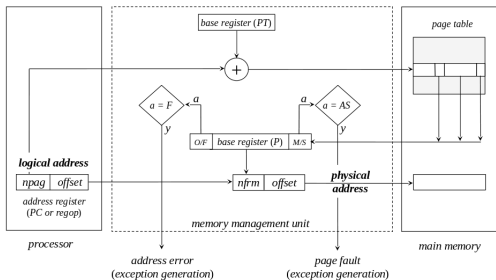
Paging

Memory management unit (MMU)



Paging


Logical address to physical address translation



- Structuring the logical address space of the process in order to map the whole, or at least a fraction, of the address space provided by the processor (in any case, always greater than or equal to the size of the existing main memory), it becomes possible to eliminate the need for the **limit register** associated with the size of the page table
- As a consequence, the gap between the heap memory and the stack can be maximized
- There is not **limit register** associated with the page

Paging

Page table entries



- Page table contains one entry per page
- Entry definition:

O/F	M/S	ref	mod	perm	frame number	block number in swap area
-----	-----	-----	-----	------	--------------	---------------------------

- **O/F** – flag indicating if page has been already assigned to process
- **M/S** – flag indicating if page is in memory
- **ref** – flag indicating if page has been referenced
- **mod** – flag indicating if page has been modified
- **perm** – permissions
- **frame number** – frame where page is, if in memory
- **block number in swap area** – block where page is, in swapping area

- Advantages:

- **general** – the scope of your application is independent of the type of processes that will be executed (number and size of their address spaces)
- **good usage of main memory** – does not lead to external fragmentation and internal fragmentation is practically negligible
- **does not have special hardware requirements** – the memory management units in today's general-purpose processors implements it

- Disadvantages:

- **longer memory access** – double access to memory, because of a prior access to the page table
 - Existence of a TLB (translation lookaside buffer) minimizes the impact
- **very demanding operability** – requires the existence of a set of support operations, that are complex and have to be carefully designed to not compromise efficiency

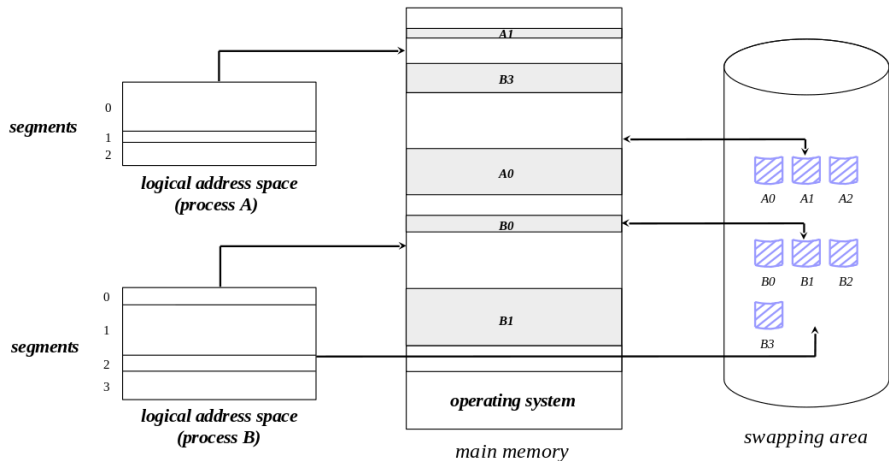
Segmentation

Introduction

- Typically, the logical address space of a process is composed of different type of segments:
 - **code** – one segment per code module
 - **static variables** – one segment per module containing static variables
 - **heap memory** – one segment
 - **shared memory** – one segment per shared region
 - **stack** – one segment
 - Different segments may have different sizes
- In a **segmentation architecture**, the segments of a process are manipulated separately
 - **Dynamic partitioning** may be used to allocate each segment
 - As a consequence, a process may not be contiguous in memory
 - Even, some segments may not be in main memory

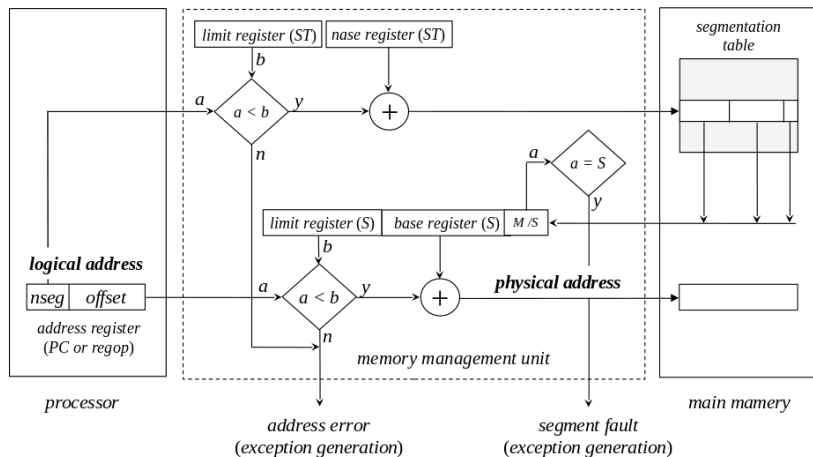
Segmentation

Illustration example



Segmentation

Memory management unit (MMU)



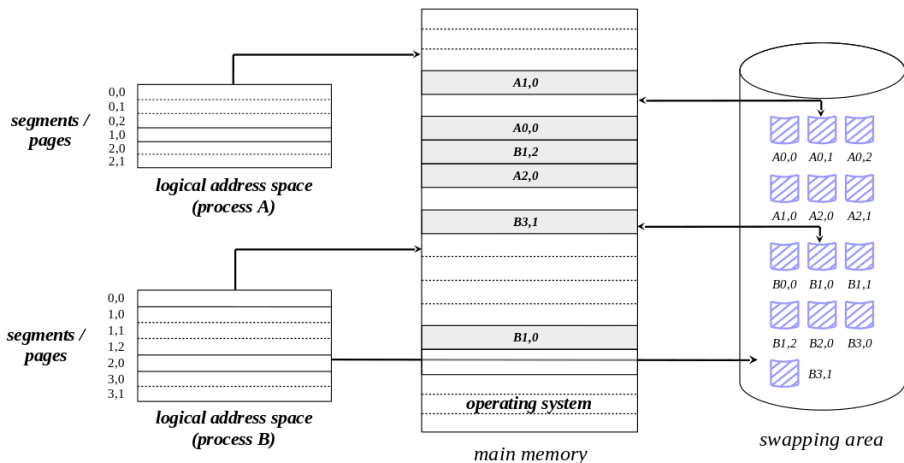
Combining segmentation and paging

Introduction

- Segmentation, taken alone, can have some drawbacks:
 - It may result in external fragmentation
 - A growing segment can impose a change in its location
- Merging segmentation and paging can solve these issues
 - First, the logical address space of a process is partitioned into segments
 - Then, each segment is divided into pages
- However, this introduces a growing complexity

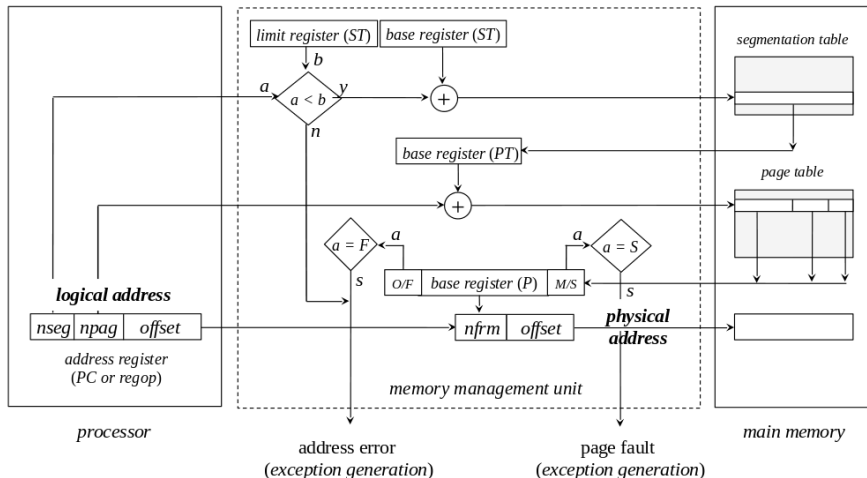
Combining segmentation and paging

Illustration example




Combining segmentation and paging

Memory management unit (MMU)



Combining segmentation and paging

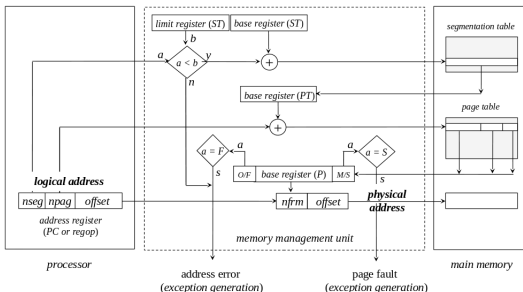
Logical address to physical address translation



- The MMU must contain 3 base registers and 1 limit register
 - 1 base register for the segmentation table
 - 1 limit register for the segmentation table
 - 1 base register for the page table
 - 1 base register for the memory frame

Combining segmentation and paging

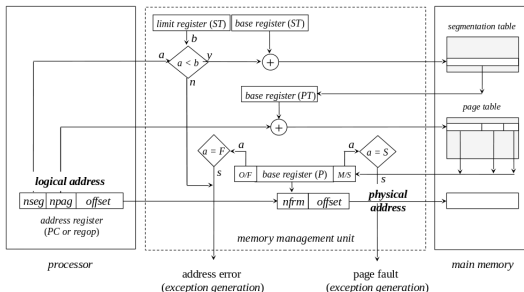
Logical address to physical address translation (2)



- An access to memory unfolds into 3 steps:
 - Access to the segmentation table
 - Access to the page table
 - Access to the physical address

Combining segmentation and paging

Logical address to physical address translation (3)



- Entry of the segmentation table

perm	memory address of the page table
------	----------------------------------


- Entry of the page table

O/F	M/S	ref	mod	frame number	block number in swap area
-----	-----	-----	-----	--------------	---------------------------

- The `perm` field is now associated to the segment

Page replacement

Introduction




- In a paging (or combination of segmentation and paging) architecture, memory is partitioned into frames, each the same size as a page
 - A frame may be either **free** or **occupied** (containing a page)
- A page in memory may be:
 - **locked** – if it can not be removed from memory (kernel, buffer cache, memory-mapped file)
 - **unlocked** – if it can be removed from memory
- If no free frame is available, a occupied one may need to be released
 - This is the purpose of **page replacement**
- Page replacement only applies to unlocked pages

Page replacement


Lists of free and occupied frames

- Free frames are organized in a list – **list of free frames**
- Occupied frames, associated to unlocked pages, are also organized in a list – **list of occupied frames**
- How the list of occupied frames is organized depends on the **page replacement policy**

list of busy frames



list of free frames



Page replacement

Action on page fault

- On **page fault**, if the list of free frames is empty, a occupied frame must be selected for replacement
- Alternatively, the system can promote page replacement as to maintain the list of free frames always with some elements
 - This allows to load the faulty page and free a occupied frame at the same time
- The question is: which frame should be selected for replacement?
- An **optimal policy** selects for replacement that page for which the time to the next reference is the longest
 - Unless we have a Crystal-Ball, it is **impossible to implement**
 - But, useful as benchmark
- Covered algorithms:
 - **Least Recently Used (LRU)**
 - **Not Recently Used (NRU)**
 - **First In First Out (FIFO)**
 - **Second chance**
 - **Clock**

Page replacement policies

LRU algorithm

- The **Least Recently Used** policy selects for replacement the frame that has not been referenced the longest
 - Based on the principle of locality of reference, if a frame is not referenced for a long time, it is likely that it will not be referenced in the near future
- Each frame must be labelled with the time of the last reference
 - Additional specific hardware may be required
- On page replacement, the list of occupied frames must be traversed to find out the one with the oldest last access time
- High cost of implementation and not very efficient

Page replacement policies

NRU algorithm

- The **Not Recently Used** policy selects for replacement a frame based on classes
- Bits **Ref** and **Mod**, fields of the entry of the page table, and typically processed by conventional MMU, are used to define classes of frames

class	Ref	Mod
0	0	0
1	0	1
2	1	0
3	1	1

- On page replacement, the algorithm selects at random a frame from the lowest non-empty class
- Periodically, the system traverse the list of occupied frames and put **Ref** at zero

Page replacement policies

FIFO algorithm

- The **FIFO** policy selects for replacement based on the length of stay in memory
 - Based on the assumption that the longer a page resides in memory, the less likely it is to be referenced in the future
- The list of occupied frames is considered to be organized in a FIFO that reflects the loading order of the corresponding pages in main memory
- On page replacement, the frame with the oldest page is selected
- The assumption in itself is extremely fallible
 - Consider for instance system shared libraries
 - But can be interesting with a refinement

Page replacement policies


Second chance algorithm

- The **second chance** policy is an improvement of the FIFO algorithm, giving a page a second chance before it is replaced
- On page replacement:
 - The frame with the oldest page is selected as a candidate
 - If its `Ref` bit is at zero, the selection is done
 - If not, the `Ref` bit of the candidate frame is reset, the frame is inserted again in the FIFO, and the process proceeds with the next frame
 - The process ends when a frame with the `Ref` bit at zero is found
- Note that such frame is always found

Page replacement policies

Clock algorithm

- The **clock** policy is an improvement of the second chance algorithm, avoiding the removal and reinsetion of elements in the FIFO
- The list is transformed into a circular one and a pointer signals the oldest element
 - The action of removal followed by a reinsertion corresponds to a pointer advance
- On page replacement:
 - While the `Ref` bit of a frame is non-zero, that bit is reset and the pointer advances to the next frame
 - The first frame with the `Ref` bit at zero is chosen for replacement
 - After replacement, the pointer is placed pointing to the next element



Working set

- Assume that initially only 2 pages of a process are in memory
 - The one containing the first instruction
 - The one containing the start of the stack
- After execution starts and for a while, page faults will be frequent
- Then the process will enter a phase in which page faults will be almost inexistent
 - Corresponds to a period where, accordingly to the principle of locality of reference, the fraction of the address space that the process is currently referencing is all present in main memory
- This set of pages is called the **working set** of the process
- Over time the working set of the process will vary, not only with respect to the number, but also with the specific pages that define it

Thrashing

- Consider that the maximum number of frames assigned to a process is fixed
- If this number is always greater or equal to the number of pages of the different working sets of the process:
 - the process's life will be a succession of periods with frequent page faults with periods almost without them
- If it is lower
 - the process will be continuously generating page faults
 - in such cases, it is said to be in [thrashing](#)
- Keeping the working set of a process always in memory is page replacement design challenge

Demand paging vs. prepaging

- When a process transition to the ready state, what pages should be placed in main memory?
- Two possible strategies: **demand paging** and **prepaging**
- **Demand paging** – place none and wait for the page faults
 - inefficient
- **Prepaging** – place those most likely to be referenced
 - first time, the two pages mentioned before (code and stack)
 - next times, those that were in main memory when the process was suspended
 - more efficient

Bibliography

- Operating Systems: Internals and Design Principles, W. Stallings, Prentice-Hall International Editions, 7th Ed, 2012
 - Chapter 7: Memory Management (sections 7.1 to 7.4)
 - Chapter 8: Virtual Memory (sections 8.1 to 8.2)
- Operating Systems Concepts, A. Silberschatz, P. Galvin and G. Gagne, John Wiley & Sons, 9th Ed, 2013
 - Chapter 8: Main Memory (sections 8.1 to 8.5)
 - Chapter 9: Virtual Memory (sections 9.1 to 9.6)
- Modern Operating Systems, A. Tanenbaum and H. Bos, Pearson Education Limited, 4th Ed, 2015
 - Chapter 3: Memory Management (section 3.1 to 3.7)